

Hardware-Competitive Quantum Reservoir Computing for SPY Volatility

A native-hardware hybrid across IQM Emerald, QuEra Aquila, and cloud Ising machines

Track A: Financial Volatility Prediction | Team Quanties | Draft 15 July 2026

Abstract

We forecast next-day annualized 21-day realized volatility of SPY from public daily OHLCV data (2010-2024). GARCH(1,1) supplies a structural prior; a fixed quantum reservoir learns only a regularized correction to the GARCH log-volatility forecast. Zero correction therefore equals GARCH exactly, giving the quantum component a stringent role: extract nonlinear residual structure without replacing a strong econometric model.

The primary real-QPU result is a 9-qubit native 3x3 IQM Emerald reservoir evaluated on 120 chronological days at 500 shots per circuit. Its pre-specified QREM plus affine-transfer path gives RMSE 0.00831368 versus 0.00832112 for GARCH, a 0.0895% point improvement. The moving-block-bootstrap 95% interval for the RMSE difference is $[-0.0000608, +0.0000262]$, so we claim hardware-level competitiveness, not statistically established quantum advantage. QLIKE is 0.016538 versus 0.016476 for GARCH; the Mincer-Zarnowitz joint p-value is 0.166, so forecast unbiasedness is not rejected.

Headline. Across executed substrates, the hybrid remains within 0.43% RMSE of matched GARCH. Real Aquila transfers from its local AHS simulator within 0.029% RMSE, while a 12-qubit IQM expansion improves feature fidelity but worsens forecasting. Native topology and validation-locked regularization matter more than width alone.


1. Problem framing and data

The target $RV[t]$ is the rolling 21-day standard deviation of log returns multiplied by $\sqrt{252}$. At forecast time, 20 of the 21 squared-return terms are known, giving GARCH a structural advantage. Inputs are causal lags of returns, absolute returns, squared returns, and heterogeneous autoregressive realized-volatility summaries. All splits are chronological; model selection uses training-only expanding folds.

We report RMSE, volatility-specific QLIKE loss, and Mincer-Zarnowitz forecast calibration. Comparators include Persistence, Ridge/AR, ESN-200, LSTM, and GARCH(1,1). A 746-day full benchmark establishes classical/simulator performance; a separate 400-train/120-test window supports hardware-ready and real-hardware transfer experiments.

2. Hybrid QRC architecture

For each market state $x[t]$, GARCH produces a log-volatility prior. A fixed reservoir maps $x[t]$ to shot-estimated one- and two-body observables; only a Ridge readout is trained. The forecast is $\log RV_{\text{hat}}[t] = \log RV_{\text{GARCH}}[t] + c f_{\text{QRC}}(x[t])$, with the Ridge penalty and correction strength c locked on validation data. Quantum parameters are fixed before the test window; no variational quantum optimization is used.



Train only the classical head; keep the physical reservoir fixed and evaluate chronologically.

Figure 1. The GARCH prior protects baseline performance; the fixed reservoir supplies nonlinear residual features. Hardware mitigation is label-free.

Substrate	Native encoding and dynamics	Readout	Purpose
IQM Emerald	Rz input on a selected nearest-neighbor grid; fixed CZ mixing	45 Z/ZZ	Primary real gate-QPU forecast
QuEra Aquila	Negative local detuning on a valid irregular atom grid; Rydberg AHS	5 Z + 10 ZZ	Real analog-QPU transfer
Amplify AE / Toshiba	Input fields h_i with fixed signed all-to-all J_{ij} ; cloud optimization/sampling	10 Z + 45 ZZ	Classical signed-Ising ablation
TFIM simulator	Signed J_{ij} plus transverse-field evolution	Z/ZZ ensemble	Mechanism and advantage candidate

Why a reservoir is appropriate

Volatility is nonlinear, heteroskedastic, and regime switching. A fixed interacting system supplies a high-dimensional nonlinear projection and fading memory, while the linear head controls variance. The residual formulation tests whether the reservoir adds orthogonal structure after the known 21-day decomposition is supplied. Z and ZZ observables expose local response and interaction-mediated correlations.

Hardware transfer is conservative. IQM circuits use explicit native qubit selection and reject SWAPs; Aquila geometry satisfies the 4 micrometer per-axis rule and uses local detuning. QREM and affine feature alignment use calibration information or unlabeled features only, preventing target leakage.

3. Experimental design and simulator evidence

The 746-day benchmark establishes task difficulty. GARCH leads the raw-task baselines (RMSE 0.007948), followed by Persistence (0.009408), ESN-200 (0.009704), and LSTM (0.010662). On the GARCH-residual target, the 10-qubit transverse-field Ising reservoir reaches RMSE 0.007844 and QLIKE 0.006445 versus 0.006863 for GARCH: improvements of 1.31% and 6.09%. The identical Ridge residual ablation has RMSE 0.008051. This is a simulator result, not an executed hardware-advantage claim.

Simulator evidence: advantage candidate and hardware-ready ablations




Figure 2. Left: the simulated TFIM hybrid is the advantage candidate. Right: every hardware-ready reservoir improves on the same linear residual ablation, although none beats GARCH on the common 120-day window.

Selection, noise, and ablation controls

- All hyperparameters, layouts, Ridge penalties, and correction strengths are selected on expanding training/validation folds. Test rows are scored once.
- The final 9-qubit configuration gains 0.536% over GARCH in exact statevector simulation and 0.410% under the IQM-matched 500-shot proxy.
- On the common 400/120 window, Pasqal analog, gate QRC, QuEra analog, and signed-Ising reservoirs improve over the linear residual ablation by 2.36%, 2.10%, 1.79%, and 1.52%.
- The 12-qubit native-grid check raises hardware-to-exact feature correlation from 0.770 to 0.818 but changes the forecast from 0.090% better than GARCH to 0.196% worse.

Mechanism interpretation

Removing the transverse field leaves the cloud signed-Ising tier near GARCH but does not reproduce the full simulated gain; removing signed programmability produces the neutral-atom tier with the same conclusion. The evidence supports a falsifiable hypothesis - useful residual features arise from transverse-field mixing plus signed couplings - while real hardware supports competitiveness and portability rather than advantage.

4. Executed hardware and cloud results

All rows below are completed external executions and are compared with GARCH on identical held-out rows. Percentages are comparable within a row even when absolute RMSE differs across windows.

Real execution: near-GARCH performance and a topology-aware scaling result




Figure 3. Negative RMSE gap is better than GARCH. The small 9-qubit point gain is not significant; width alone does not improve the forecast.

Platform	Native workload	Evaluation	Cost	Gap	Reading
IQM 9q	depth 32; 24 CZ; 0 SWAP	120 x 500	20.25 Resonance cr	-0.090%	Primary QPU; CI crosses zero
IQM 12q	depth 38; 34 CZ; 0 SWAP	120 x 500	20.25 Resonance cr	+0.196%	Scaling negative result
Aquila 5 atoms	native AHS; local detuning	23 x 50; 20 scored	1,840 qBraid cr	+0.182%	Only 0.029% behind local AHS
Amplify AE	10 spins; 45 signed J	120 solves	existing credits	+0.055%	Classical ablation
Toshiba	10 spins; 45 signed J	120 x 200 states	existing credits	+0.421%	Classical ablation

Primary IQM result

Emerald job 019f5691-063b-7a11-8bbc-db5057b79259 executed 60,000 shots on QB17-19, QB25-27, and QB33-35. QREM plus feature-only affine transfer gives feature correlation 0.770 and RMSE 0.00831368. The bootstrap probability of beating GARCH is 0.658; the interval crosses zero. QLIKE is slightly worse, so the defensible conclusion is GARCH-level real-QPU performance with a small positive RMSE point estimate.

Analog and cloud-Ising transfer

Aquila completed 23 paid tasks after zero-cost validation exposed native geometry and local-detuning constraints. Three unlabeled rows calibrate feature alignment and 20 are scored. Hardware RMSE is 0.0137535 versus 0.0137495 locally and 0.0137285 for GARCH; it improves 6.06% over the Ridge residual ablation. Amplify AE and Toshiba remain within 0.43% of GARCH but are classical optimizers, not quantum-advantage evidence.

5. Conclusions, impact, and reproducibility

Conclusions

- The strongest executed claim is cross-platform, GARCH-competitive hybrid QRC: IQM is 0.0895% better by RMSE point estimate; Aquila is 0.1823% behind GARCH and effectively matches local AHS; all executed tiers stay within 0.43%.
- The strongest mechanism claim is simulated: the transverse-field signed-coupling reservoir improves RMSE and QLIKE over GARCH, whereas neutral-atom and classical signed-Ising ablations do not reproduce the full gain.
- More qubits are not automatically better. Native connectivity, noise-aware training, low depth, and conservative correction strength dominate width in this data regime.

Stakeholder relevance

Daily volatility forecasts support hedging, market-maker inventory and spread decisions, portfolio risk limits, and scenario generation. The deployable contribution is a modular residual feature engine: it can be switched off to recover GARCH exactly, deployed across substrates, and monitored through correction magnitude and calibration metrics.

Limitations

- The IQM RMSE gain is small and not significant; QLIKE does not improve. Aquila has only 20 scored days and 50 shots per task.
- Amplify AE and Toshiba lack coherent transverse-field dynamics. The GARCH-beating TFIM result is simulated, not QPU-executed.
- The exact hardware window still needs GJR-GARCH, EGARCH, and HAR-RV; the full-window panel already includes GARCH, ESN, LSTM, Ridge/AR, and Persistence.

Reproducibility and submission status

Requirement	Evidence	Status
Concrete qBraid execution	Aquila task IDs/checkpoint; qir-sv records	Complete
Re-runnable workflow	Judge notebook, experiments/, compact CSV/JSON evidence	Complete
5-page write-up	This body plus separate references	Draft complete
Cover + Launch button	README Launch button added; current cover is Phase 2	Cover pending

Recommended final work

- 1 Do not spend the remaining 580 qBraid credits. If 2,500 credits are added, authorize Aquila only after a validation-locked local analog candidate beats GARCH by at least 0.30% RMSE.
- 2 If that gate passes, use 31 x 50-shot tasks (estimated 2,480 credits): four engineering-smoke tasks, then resume the identical checkpoint; reserve three rows for label-free transfer and score 28.
- 3 Prioritize free work first: exact-window GJR/EGARCH/HAR baselines, regime/bootstrap analysis, the qBraid notebook, and a deterministic data snapshot.
- 4 No further paid IQM, Amplify, or Toshiba run is justified unless the rubric changes; current runs already answer the main scaling and substrate questions.

Draft completion gate. Insert the official Phase 3 cover without altering its template, replace the Phase 2 cover, add the final commit hash, verify the qBraid Launch link, then remove this box before submission.

References (excluded from the five-page limit)

- [1] qBraid, MITRE, and JonesTrading. *Global Industry Challenge 2026: Phase 3 Challenge Description*. 2026.
- [2] Aqora. *qBraid, MITRE & JonesTrading: GIC 2026 track page*.
aqora.io/challenges/global-industry-challenge-2026/tracks/gic-2026-qBraid-MITRE-JonesTrading
(accessed 15 July 2026).
- [3] Bollerslev, T. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics* 31, 307-327 (1986).
- [4] Corsi, F. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics* 7, 174-196 (2009).
- [5] Patton, A. J. Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics* 160, 246-256 (2011).
- [6] Jaeger, H. The echo state approach to analysing and training recurrent neural networks. GMD Report 148 (2001).
- [7] Fujii, K. and Nakajima, K. Harnessing disordered-ensemble quantum dynamics for machine learning. *Physical Review Applied* 8, 024030 (2017).
- [8] Martinez-Pena, R. et al. Dynamical phase transitions in quantum reservoir computing. *Physical Review Letters* 127, 100502 (2021).
- [9] Kornjaca, M. et al. Large-scale quantum reservoir learning with an analog quantum computer. arXiv:2407.02553 (2024).
- [10] Li, Q. et al. Quantum reservoir computing for realized volatility forecasting. arXiv:2505.13933 (2025).
- [11] Ahmed, O., Tennie, F., and Magri, L. Robust quantum reservoir computers for forecasting chaotic dynamics. *Proceedings of the Royal Society A* 481, 20250550 (2025).
- [12] qBraid. *Launch on qBraid badge and repository redirect parameters*.
github.com/qBraid/community/discussions/3 (accessed 15 July 2026).

Artifact disclosure

Generative AI tools were used for code support, debugging, figure generation, and writing assistance. The experimental design, platform choices, executed jobs, result interpretation, and final claims are the team's own. Secrets and access tokens are excluded from the repository and should be rotated after the hardware campaign.